

Convolutional Neural Networks with Alternately Updated Clique



Yibo Yang;




Zhisheng Zhong;



Tiancheng Shen;





Zhouchen Lin




Background

Deeper and deeper

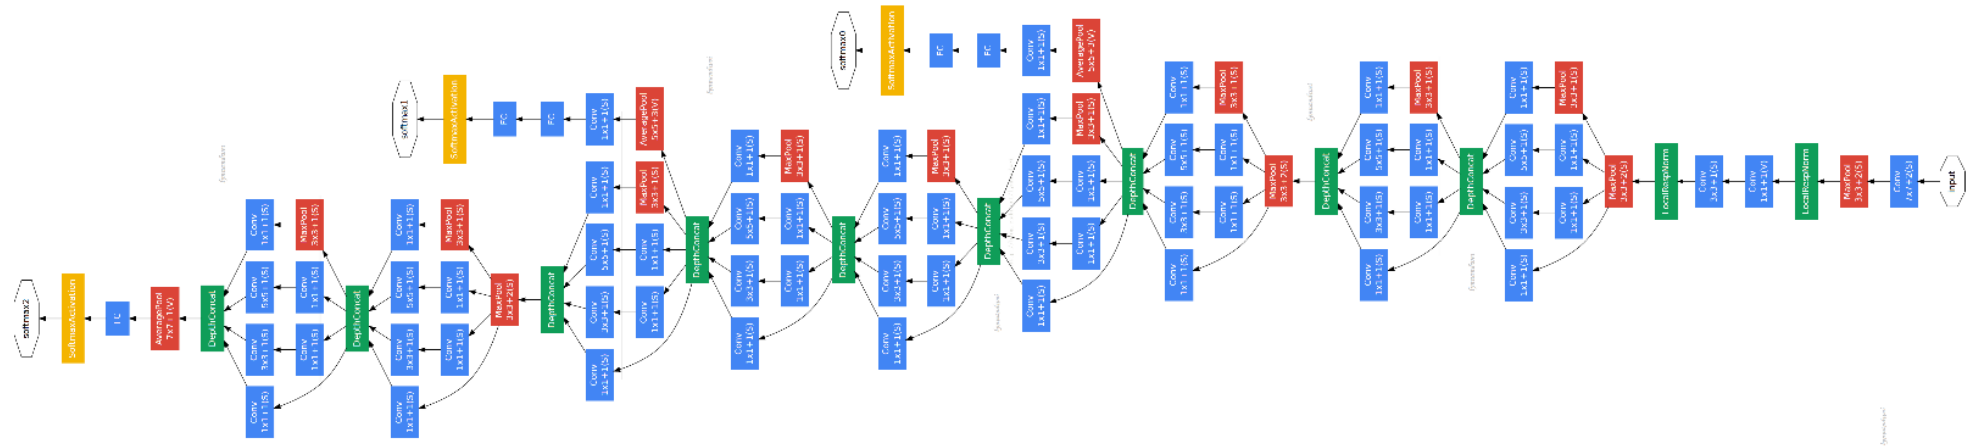


Background





A horizontal timeline with a red arrow pointing right. It marks the years 2012, 2014, 2014, 2015, and 2016. Below each year is a blue dot on the timeline line. Under each dot is the name of a neural network and its number of layers. The 2014 entry is highlighted in red.

Year	Model	Layers
2012	AlexNet	7 layers
2014	VGGNet	16/19 layers
2014	GoogLeNet	22 layers
2015	HighwayNet	100+ layers
2016	ResNet	152 layers



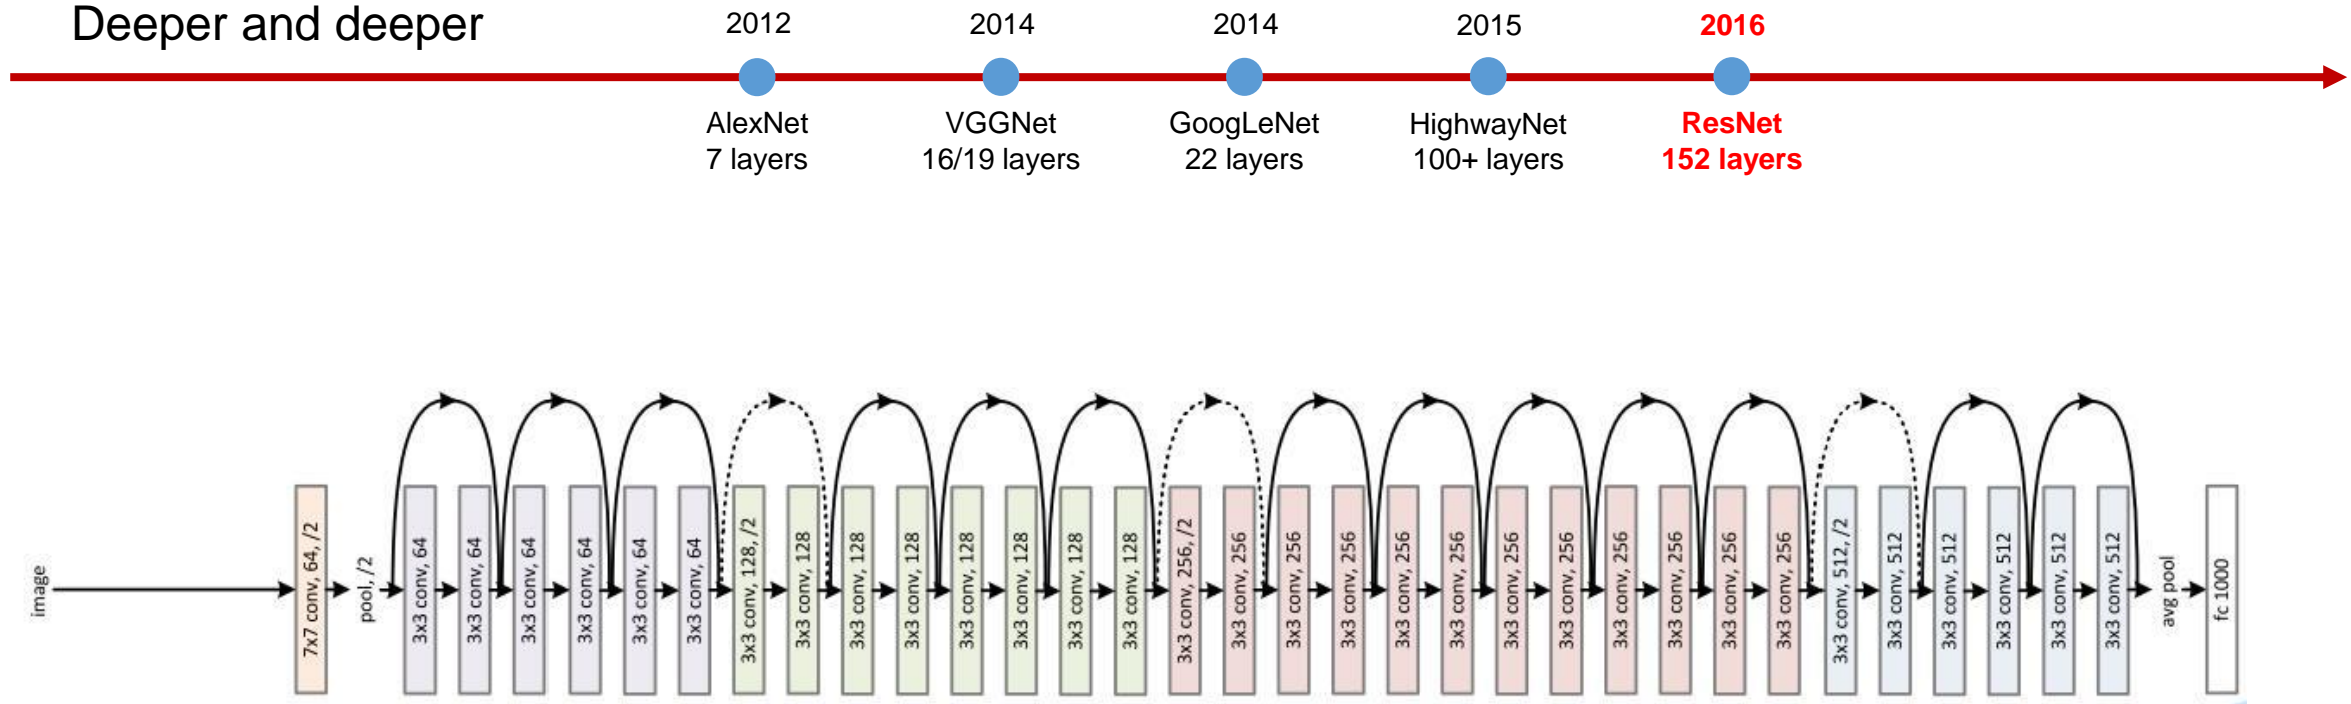
Background

Deeper and deeper



Background

Deeper and deeper



Skip Connections

Background



DenseNet, 2017

- Propagation

ResNet: $x_l = H_l(x_{l-1}) + x_{l-1}$

DenseNet: $x_l = H_l([x_0, x_1, \dots, x_{l-1}])$

- Bottleneck: BN-ReLU-Conv(1x1)-BN-ReLU-Conv(3x3)
- Compression: Reduce the number of feature maps in transition layers.
- Performance: State-of-the-art on CIFAR, SVHN datasets; parameter-efficient than ResNets.



Motivations



- **How to further maximize information flow?**



Motivations

- How to further maximize information flow?
- Does feedback connection help?



Motivations

- Attention mechanism
 - Formulating an optimization problem (Cao et al. 2015)
 - Weighting the activations spatially or channel-wisely (Chen et al. 2017, Hu et al. 2017)
 - Feedback connections (Stollenga et al. 2014, Wang et al. 2014, Zamir et al. 2017)




Motivations

- Attention mechanism
 - Formulating an optimization problem (Cao et al. 2015)
 - Weighting the activations spatially or channel-wisely (Chen et al. 2017, Hu et al. 2017)
 - **Feedback connections** (Stollenga et al. 2014, Wang et al. 2014, Zamir et al. 2017)
- Dense connections & Feedback connections
 - Bi-directionally connected -> Enable feature refinement at each update
 - Densely connected -> Information flow further maximized

Architectures

- Basic block
 - Any two layers have both forward and backward connections, and are both input and output of each other
 - More densely connected than DenseNets



Architectures

- Propagation

- First stage:

Initialize layers using feedforward dense connections

Bottom Layers	Weights	Top Layers	Feature
X_0	W_{01}	$X_1^{(1)}$	Stage-I
$\{X_0, X_1^{(1)}\}$	$\{W_{02}, W_{12}\}$	$X_2^{(1)}$	
$\{X_0, X_1^{(1)}, X_2^{(1)}\}$	$\{W_{03}, W_{13}, W_{23}\}$	$X_3^{(1)}$	
$\{X_0, X_1^{(1)}, X_2^{(1)}, X_3^{(1)}\}$	$\{W_{04}, W_{14}, W_{24}, W_{34}\}$	$X_4^{(1)}$	
$\{X_0, X_1^{(1)}, X_2^{(1)}, X_3^{(1)}, X_4^{(1)}\}$	$\{W_{05}, W_{15}, W_{25}, W_{35}, W_{45}\}$	$X_5^{(1)}$	
$\{X_2^{(1)}, X_3^{(1)}, X_4^{(1)}, X_5^{(1)}\}$	$\{W_{21}, W_{31}, W_{41}, W_{51}\}$	$X_1^{(2)}$	Stage-II
$\{X_3^{(1)}, X_4^{(1)}, X_5^{(1)}, X_1^{(2)}\}$	$\{W_{32}, W_{42}, W_{52}, W_{12}\}$	$X_2^{(2)}$	
$\{X_4^{(1)}, X_5^{(1)}, X_1^{(2)}, X_2^{(2)}\}$	$\{W_{43}, W_{53}, W_{13}, W_{23}\}$	$X_3^{(2)}$	
$\{X_5^{(1)}, X_1^{(2)}, X_2^{(2)}, X_3^{(2)}\}$	$\{W_{54}, W_{14}, W_{24}, W_{34}\}$	$X_4^{(2)}$	
$\{X_1^{(2)}, X_2^{(2)}, X_3^{(2)}, X_4^{(2)}\}$	$\{W_{15}, W_{25}, W_{35}, W_{45}\}$	$X_5^{(2)}$	

Architectures

- Propagation

- First stage:

Initialize layers using feedforward dense connections

- Second stage:

Alternate updating rule

$$X_i^{(k)} = g \left(\sum_{l < i} W_{li} * X_l^{(k)} + \sum_{m > i} W_{mi} * X_m^{(k)} \right)$$

($i \geq 1, k \geq 2, g$: non-linear activation)


Bottom Layers	Weights	Top Layers	Feature
X_0	W_{01}	$X_1^{(1)}$	Stage-I
$\{X_0, X_1^{(1)}\}$	$\{W_{02}, W_{12}\}$	$X_2^{(1)}$	
$\{X_0, X_1^{(1)}, X_2^{(1)}\}$	$\{W_{03}, W_{13}, W_{23}\}$	$X_3^{(1)}$	
$\{X_0, X_1^{(1)}, X_2^{(1)}, X_3^{(1)}\}$	$\{W_{04}, W_{14}, W_{24}, W_{34}\}$	$X_4^{(1)}$	
$\{X_0, X_1^{(1)}, X_2^{(1)}, X_3^{(1)}, X_4^{(1)}\}$	$\{W_{05}, W_{15}, W_{25}, W_{35}, W_{45}\}$	$X_5^{(1)}$	
$\{X_2^{(1)}, X_3^{(1)}, X_4^{(1)}, X_5^{(1)}\}$	$\{W_{21}, W_{31}, W_{41}, W_{51}\}$	$X_1^{(2)}$	Stage-II
$\{X_3^{(1)}, X_4^{(1)}, X_5^{(1)}, X_1^{(2)}\}$	$\{W_{32}, W_{42}, W_{52}, W_{12}\}$	$X_2^{(2)}$	
$\{X_4^{(1)}, X_5^{(1)}, X_1^{(2)}, X_2^{(2)}\}$	$\{W_{43}, W_{53}, W_{13}, W_{23}\}$	$X_3^{(2)}$	
$\{X_5^{(1)}, X_1^{(2)}, X_2^{(2)}, X_3^{(2)}\}$	$\{W_{54}, W_{14}, W_{24}, W_{34}\}$	$X_4^{(2)}$	
$\{X_1^{(2)}, X_2^{(2)}, X_3^{(2)}, X_4^{(2)}\}$	$\{W_{15}, W_{25}, W_{35}, W_{45}\}$	$X_5^{(2)}$	

Architectures

- Propagation


$$X_i^{(k)} = g \left(\sum_{l < i} W_{li} * X_l^{(k)} + \sum_{m > i} W_{mi} * X_m^{(k)} \right)$$

$$X_1^{(2)} = g(W_{21} * X_2^{(1)} + W_{31} * X_3^{(1)} + W_{41} * X_4^{(1)} + W_{51} * X_5^{(1)})$$



Architectures

- Propagation




$$X_i^{(k)} = g \left(\sum_{l < i} W_{li} * X_l^{(k)} + \sum_{m > i} W_{mi} * X_m^{(k)} \right)$$

$$X_1^{(2)} = g(W_{21} * X_2^{(1)} + W_{31} * X_3^{(1)} + W_{41} * X_4^{(1)} + W_{51} * X_5^{(1)})$$

$$X_2^{(2)} = g(W_{12} * X_1^{(2)} + W_{32} * X_3^{(1)} + W_{42} * X_4^{(1)} + W_{52} * X_5^{(1)})$$

Architectures

- Propagation



$$X_i^{(k)} = g \left(\sum_{l < i} W_{li} * X_l^{(k)} + \sum_{m > i} W_{mi} * X_m^{(k)} \right)$$


$$X_1^{(2)} = g(W_{21} * X_2^{(1)} + W_{31} * X_3^{(1)} + W_{41} * X_4^{(1)} + W_{51} * X_5^{(1)})$$

$$X_2^{(2)} = g(W_{12} * X_1^{(2)} + W_{32} * X_3^{(1)} + W_{42} * X_4^{(1)} + W_{52} * X_5^{(1)})$$

$$X_3^{(2)} = g(W_{13} * X_1^{(2)} + W_{23} * X_2^{(2)} + W_{43} * X_4^{(1)} + W_{53} * X_5^{(1)})$$

Architectures

- Propagation



$$X_i^{(k)} = g \left(\sum_{l < i} W_{li} * X_l^{(k)} + \sum_{m > i} W_{mi} * X_m^{(k)} \right)$$

$$X_1^{(2)} = g(W_{21} * X_2^{(1)} + W_{31} * X_3^{(1)} + W_{41} * X_4^{(1)} + W_{51} * X_5^{(1)})$$


$$X_2^{(2)} = g(W_{12} * X_1^{(2)} + W_{32} * X_3^{(1)} + W_{42} * X_4^{(1)} + W_{52} * X_5^{(1)})$$

$$X_3^{(2)} = g(W_{13} * X_1^{(2)} + W_{23} * X_2^{(2)} + W_{43} * X_4^{(1)} + W_{53} * X_5^{(1)})$$

$$X_4^{(2)} = g(W_{14} * X_1^{(2)} + W_{24} * X_2^{(2)} + W_{34} * X_3^{(2)} + W_{54} * X_5^{(1)})$$

Architectures

- Propagation



$$X_i^{(k)} = g \left(\sum_{l < i} W_{li} * X_l^{(k)} + \sum_{m > i} W_{mi} * X_m^{(k)} \right)$$

$$X_1^{(2)} = g(W_{21} * X_2^{(1)} + W_{31} * X_3^{(1)} + W_{41} * X_4^{(1)} + W_{51} * X_5^{(1)})$$

$$X_2^{(2)} = g(W_{12} * X_1^{(2)} + W_{32} * X_3^{(1)} + W_{42} * X_4^{(1)} + W_{52} * X_5^{(1)})$$

$$X_3^{(2)} = g(W_{13} * X_1^{(2)} + W_{23} * X_2^{(2)} + W_{43} * X_4^{(1)} + W_{53} * X_5^{(1)})$$

$$X_4^{(2)} = g(W_{14} * X_1^{(2)} + W_{24} * X_2^{(2)} + W_{34} * X_3^{(2)} + W_{54} * X_5^{(1)})$$

$$X_5^{(2)} = g(W_{15} * X_1^{(2)} + W_{25} * X_2^{(2)} + W_{35} * X_3^{(2)} + W_{45} * X_4^{(2)})$$



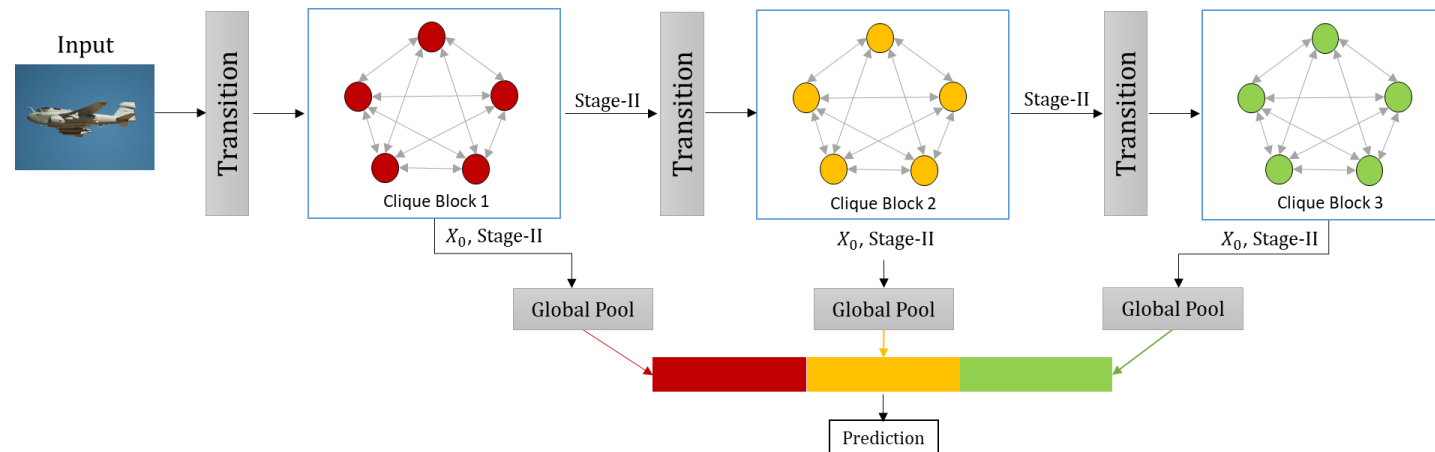
Architectures

- Propagation
 - Enable top down refinement brought by each propagation in Stage II
 - Parameters keep re-used in different stages during propagation
 - We can have Stage III and even Stage IV feature without any parameter overhead.

Architectures

- Whole structure

- A multi-scale way to compose the final representation.
- Only transit Stage-II feature into the next block to avoid progressive increment of parameters and computation.
- Multi-scale representation is more adopted in object detection or segmentation. Our multi-scale classification network is able to achieve stage-of-the-art image classification performance.





Experiments

- Ablation experiments

- CliqueNet (I+I):

Only consider **Stage-I** feature

- CliqueNet (I+II):

Use **Stage-I** feature to compose the final representation; transit **Stage-II** feature into the next block.

- CliqueNet (II+II):

Use **Stage-II** feature to compose the final representation; also transit **Stage-II** feature into the next block.

Experiments

- Ablation experiments

- CliqueNet (I+I):

Only consider **Stage-I** feature

- CliqueNet (I+II):

Use **Stage-I** feature to compose the final representation; transit **Stage-II** feature into the next block.

- CliqueNet (II+II):

Use **Stage-II** feature to compose the final representation; also transit **Stage-II** feature into the next block.

Model	Block feature	Transit	CIFAR-10
CliqueNet (I+I)	X_0 , Stage-I	Stage-I	6.64
CliqueNet (I+II)	X_0 , Stage-I	Stage-II	6.1
CliqueNet (II+II)	X_0 , Stage-II	Stage-II	5.76

(There are 36 filters in each layer, 5 layers in each block.)

Experiments

- Ablation experiments

- CliqueNet-X

Use **Stage-II** feature, but only finish the **first X steps** in the stage II.

When $X=0$, it reduces to CliqueNet (I+I)

When $X=5$, it reduces to CliqueNet (II+II)

Model	CIFAR-10	CIFAR-100
CliqueNet (X=0)	5.83	24.79
CliqueNet (X=1)	5.63	24.65
CliqueNet (X=2)	5.54	24.37
CliqueNet (X=3)	5.41	23.75
CliqueNet (X=4)	5.20	24.04
CliqueNet (X=5)	5.12	23.73

(There are 64 filters in each layer, 5 layers in each block.)



Experiments

- Experiments on CIFAR and SVHN

- Additional techniques

- Bottleneck

- Compression

- Attentional transition (SE module)

Experiments

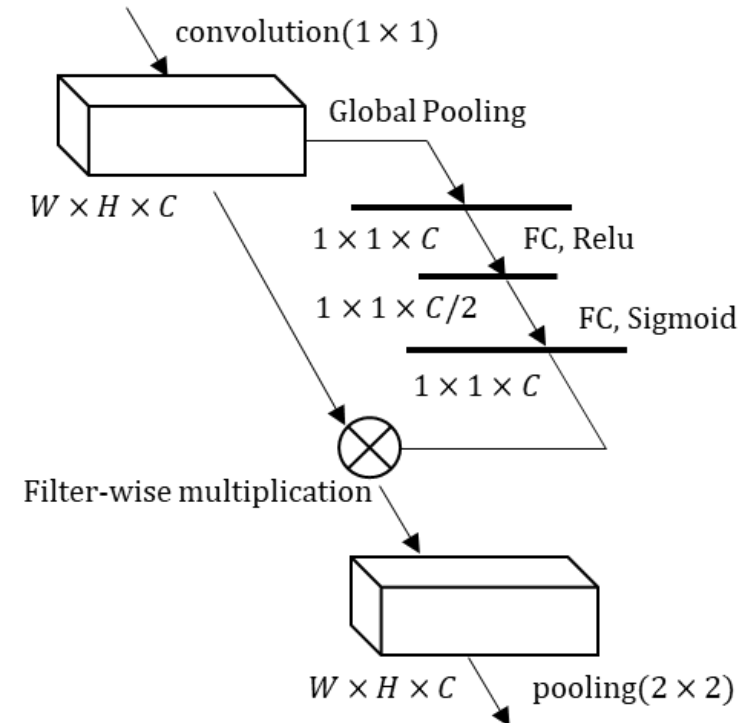
- Experiments on CIFAR and SVHN

- Additional techniques

- Bottleneck

- Compression

- Attentional transition (SE module)





Experiments

- Experiments on CIFAR and SVHN

- Additional techniques

Bottleneck

Compression

Attentional transition (SE module)

- Implementation details

300 epochs on CIFAR and 40 epochs on SVHN

0.1 learning rate initially, divided by 10 at 50% and 75% of the training procedure

Batchsize of 64 on CIFAR and SVHN

Experiments

- Experiments on CIFAR and SVHN

Model	A	B	C	FLOPs	Params	CIFAR-10	CIFAR-100	SVHN
Recurrent CNN [26]	-	-	-	-	1.86M	8.69	31.75	1.80
Stochastic Depth ResNet [18]	-	-	-	-	1.7M	11.66	37.8	1.75
dasNet [35]	-	-	-	-	-	9.22	33.78	-
FractalNet [25]	-	-	-	-	38.6M	7.33	28.2	1.87
DenseNet ($k = 12, T = 36$) [17]	-	-	-	0.53G	1.0M	7.00	27.55	1.79
DenseNet ($k = 12, T = 96$) [17]	-	-	-	3.54G	7.0M	5.77	23.79	1.67
DenseNet ($k = 24, T = 96$) [17]	-	-	-	13.78G	27.2M	5.83	23.42	1.59
CliqueNet ($k = 36, T = 12$)	-	-	-	0.91G	0.94M	5.93	27.32	1.77
CliqueNet ($k = 64, T = 15$)	-	-	-	4.21G	4.49M	5.12	23.98	1.62
CliqueNet ($k = 80, T = 15$)	-	-	-	6.45G	6.94M	5.10	23.32	1.56
CliqueNet ($k = 80, T = 18$)	-	-	-	9.45G	10.14M	5.06	23.14	1.51
DenseNet ($k = 12, T = 96$) [17]	-	✓	✓	0.58G	0.8M	5.92	24.15	1.76
DenseNet ($k = 24, T = 246$) [17]	-	✓	✓	10.84G	15.3M	5.19	19.64	1.74
CliqueNet ($k = 36, T = 12$)	✓	-	-	0.91G	0.98M	5.8	26.41	-
CliqueNet ($k = 36, T = 12$)	-	-	✓	0.98G	1.04M	5.69	26.45	-
CliqueNet ($k = 36, T = 12$)	✓	-	✓	0.98G	1.08M	5.61	25.55	1.69
CliqueNet ($k = 80, T = 15$)	✓	-	✓	6.88G	8M	5.17	22.78	1.53
CliqueNet ($k = 150, T = 30$)	✓	✓	✓	8.49G	10.02M	5.06	21.83	1.64

(Error rates on CIFAR and SVHN without any data augmentation. A, B, and C represents attentional transition, bottleneck, and compression, respectively. k is the number of filters per layer, and T is the total number of layers in three blocks.)

Experiments

- Experiments on CIFAR and SVHN

Model	A	B	C	FLOPs	Params	CIFAR-10	CIFAR-100	SVHN
Recurrent CNN [26]	-	-	-	-	1.86M	8.69	31.75	1.80
Stochastic Depth ResNet [18]	-	-	-	-	1.7M	11.66	37.8	1.75
dasNet [35]	-	-	-	-	-	9.22	33.78	-
FractalNet [25]	-	-	-	-	38.6M	7.33	28.2	1.87
DenseNet ($k = 12, T = 36$) [17]	-	-	-	0.53G	1.0M	7.00	27.55	1.79
DenseNet ($k = 12, T = 96$) [17]	-	-	-	3.54G	7.0M	5.77	23.79	1.67
DenseNet ($k = 24, T = 96$) [17]	-	-	-	13.78G	27.2M	5.83	23.42	1.59
CliqueNet ($k = 36, T = 12$)	-	-	-	0.91G	0.94M	5.93	27.32	1.77
CliqueNet ($k = 64, T = 15$)	-	-	-	4.21G	4.49M	5.12	23.98	1.62
CliqueNet ($k = 80, T = 15$)	-	-	-	6.45G	6.94M	5.10	23.32	1.56
CliqueNet ($k = 80, T = 18$)	-	-	-	9.45G	10.14M	5.06	23.14	1.51
DenseNet ($k = 12, T = 96$) [17]	-	✓	✓	0.58G	0.8M	5.92	24.15	1.76
DenseNet ($k = 24, T = 246$) [17]	-	✓	✓	10.84G	15.3M	5.19	19.64	1.74
CliqueNet ($k = 36, T = 12$)	✓	-	-	0.91G	0.98M	5.8	26.41	-
CliqueNet ($k = 36, T = 12$)	-	-	✓	0.98G	1.04M	5.69	26.45	-
CliqueNet ($k = 36, T = 12$)	✓	-	✓	0.98G	1.08M	5.61	25.55	1.69
CliqueNet ($k = 80, T = 15$)	✓	-	✓	6.88G	8M	5.17	22.78	1.53
CliqueNet ($k = 150, T = 30$)	✓	✓	✓	8.49G	10.02M	5.06	21.83	1.64

(Error rates on CIFAR and SVHN without any data augmentation. A, B, and C represents attentional transition, bottleneck, and compression, respectively. k is the number of filters per layer, and T is the total number of layers in three blocks.)

Experiments

- Experiments on ImageNet

Model	Params	Top-1	Top-5
ResNet-18	11.7M	30.43	10.76
CliqueNet-S0	5.7M	27.52	8.98
ResNet-34	21.8M	26.73	8.74
CliqueNet-S1	7.96M	26.21	8.3
DenseNet-121	7.98M	25.02	7.71
CliqueNet-S2	11M	24.82	7.51
ResNet-50	25.6M	24.01	7.02
CliqueNet-S3	14.38M	24.01	7.15


Layer	S0	S1	S2	S3
Convolution (112 × 112)	conv (7 × 7), 64, stride 2			
Pooling (56 × 56)	max pool (3 × 3), stride 2			
Block 1 (56 × 56)	36 × 5	36 × 5	36 × 5	40 × 6
Transition: conv (1 × 1), avg pool (2 × 2)				
Block 2 (28 × 28)	64 × 6	80 × 6	80 × 5	80 × 6
Transition: conv (1 × 1), avg pool (2 × 2)				
Block 3 (14 × 14)	100 × 6	120 × 6	150 × 6	160 × 6
Transition: conv (1 × 1), avg pool (2 × 2)				
Block 4 (7 × 7)	80 × 6	100 × 6	120 × 6	160 × 6

(The first number in each block is the number of filters per layer, and the second denotes the number of layers in this block)


Discussions

- Parameter efficiency

- The strong weights in DenseNets are along the diagonal.
- The strong weights in CliqueNets are distributed more evenly.
- Multi-scale structure enables our parameter efficiency.



(Our visualization)



(Visualization in DenseNet paper)

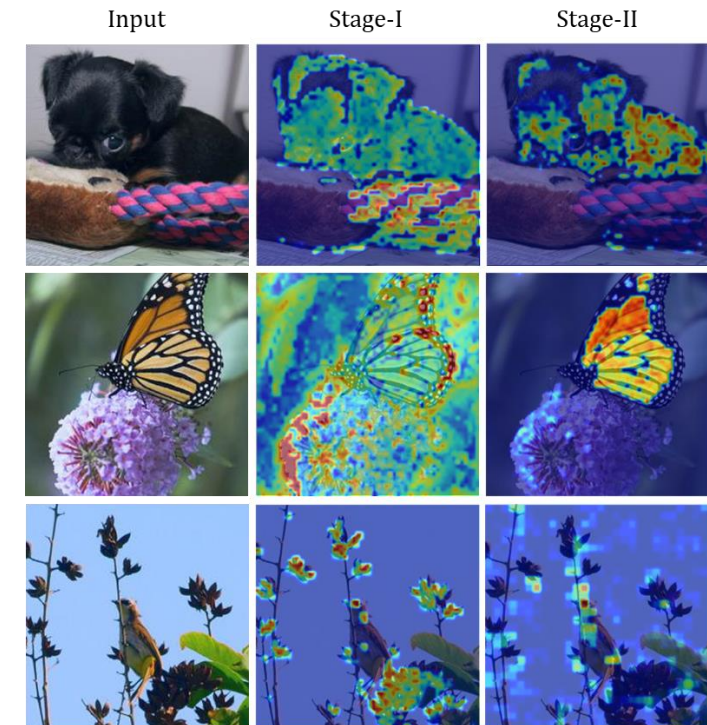
Discussions

- Parameter efficiency

- The strong weights in DenseNets are along the diagonal.
- The strong weights in CliqueNets are distributed more evenly.
- Multi-scale structure enables our parameter efficiency.

- Feature refinement

- Spatial attention is achieved to focus the activations into the target region.
- Feedback connections and alternate updating enable feature refinement.



Discussions

- Parameter efficiency

- The strong weights in DenseNets are along the diagonal.
- The strong weights in ClickNet are distributed more globally.
- Multi-scale structure enables our parameter efficiency.

Thanks for your attention!

- Feature refinement

- Spatial attention is achieved to focus the activations into the target region.
- Feedback connections and alternate updating enable feature refinement.

